Slides: https://magnusviri.com/saintcon2026

AI Coding Agents

magnusviri - James Reynolds, October 28, 2026

About magnusviri - James Reynolds

- U of U BMus 1999
- U of U Marriott Library, 14 years (Mac admin)
- U of U School of Biological Sciences/COS/LAS, 12 years (Mac, Linux, network admin)
- Security has been a “hobby” of mine for decades
- Since September 2022 (Stable Diffusion 1.0) I’ve been “obsessed with AI”
- https://magnusviri.com/saintcon2026

Agenda

- Demo
- Why Learn to use AI Agents?
- The Basics
- What Could Possibly Go Wrong?
- Mostly Secure Setup
- Coding Tips
- Demo Wrap Up

Demo

Every Demo Takes Time

- Ironicly, typing a prompt and then doing something else is the new normal
- Let’s start demoing NOW!
- tart is a macOS virtualization app
  - Was Cirruslabs but OpenAI bought it [seriously?]
  - Let’s add an MCP server

Let’s Plan

- Goal: an AI agent can autonomously control and ssh into macOS VMs
- The server will run on the macOS host
- The client will run in a Docker container
- Needs to run tart clone, run, ip, and stop
- Needs to be able to ssh into VM (but none others?)
- Needs to be able to check host disk space

Why Learn to use AI Agents?

Open Source Summit, May 2026

- https://www.youtube.com/watch?v=fi29pfLcW4I

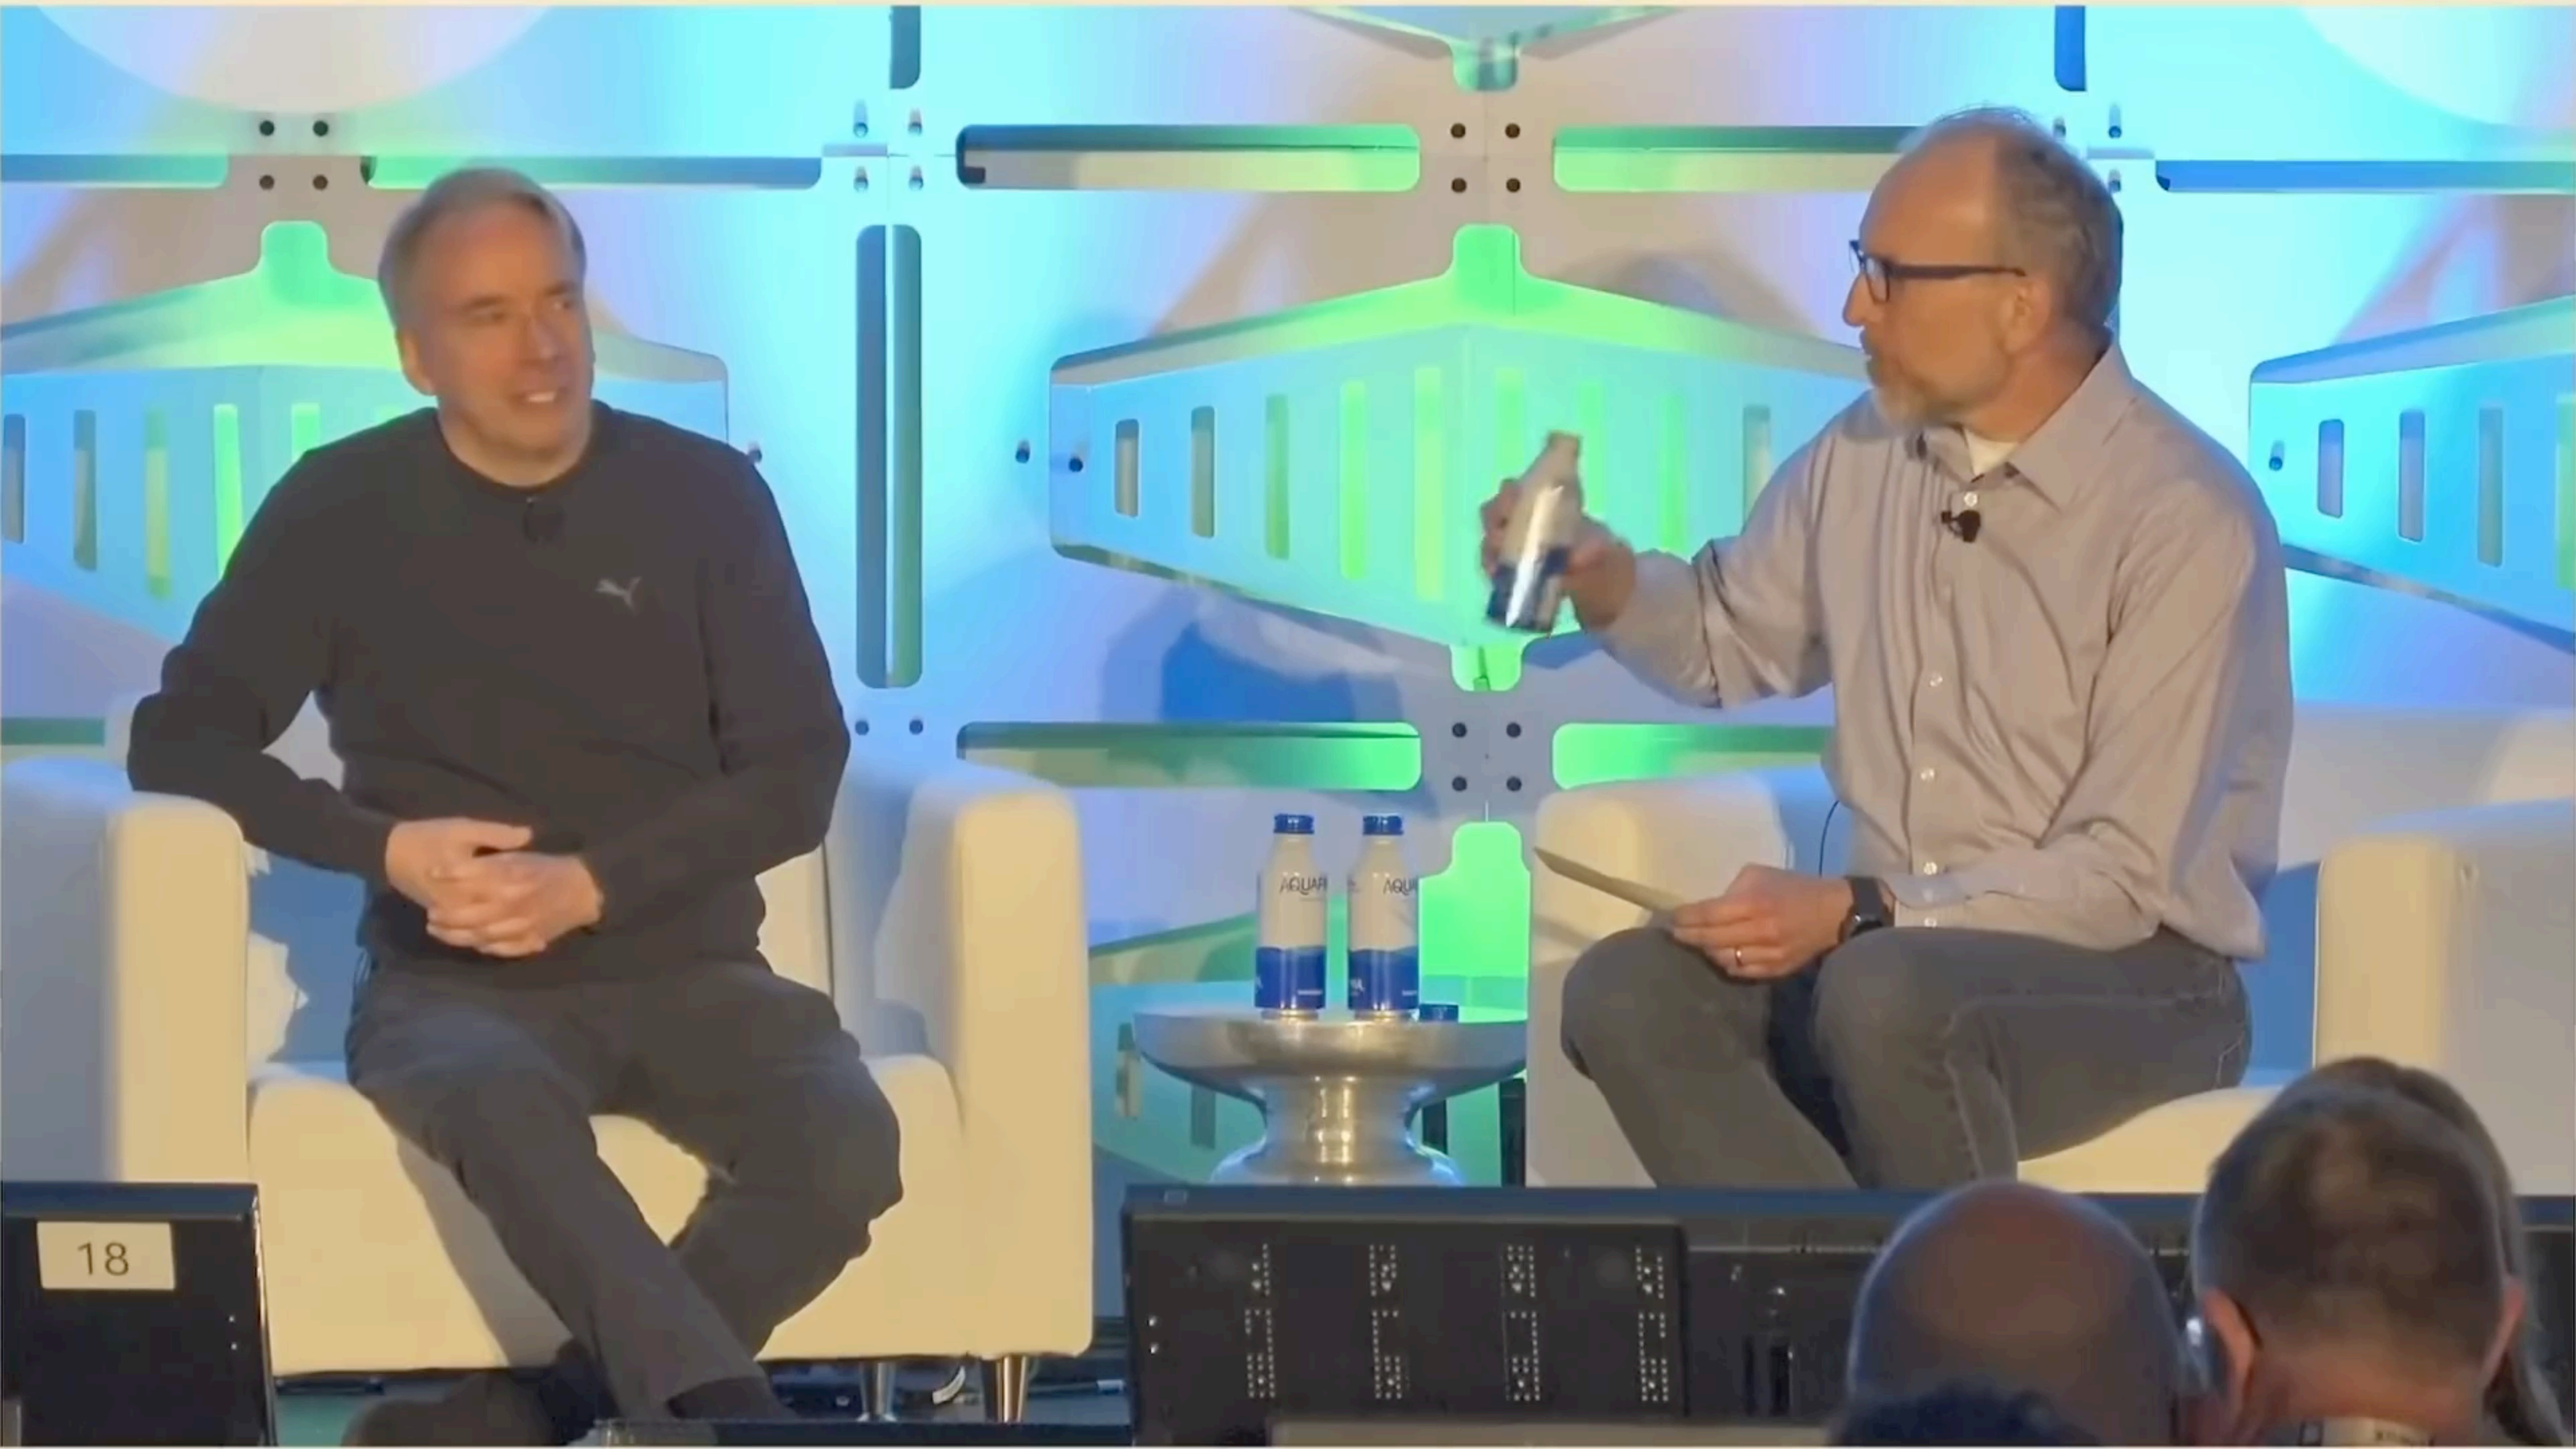
18

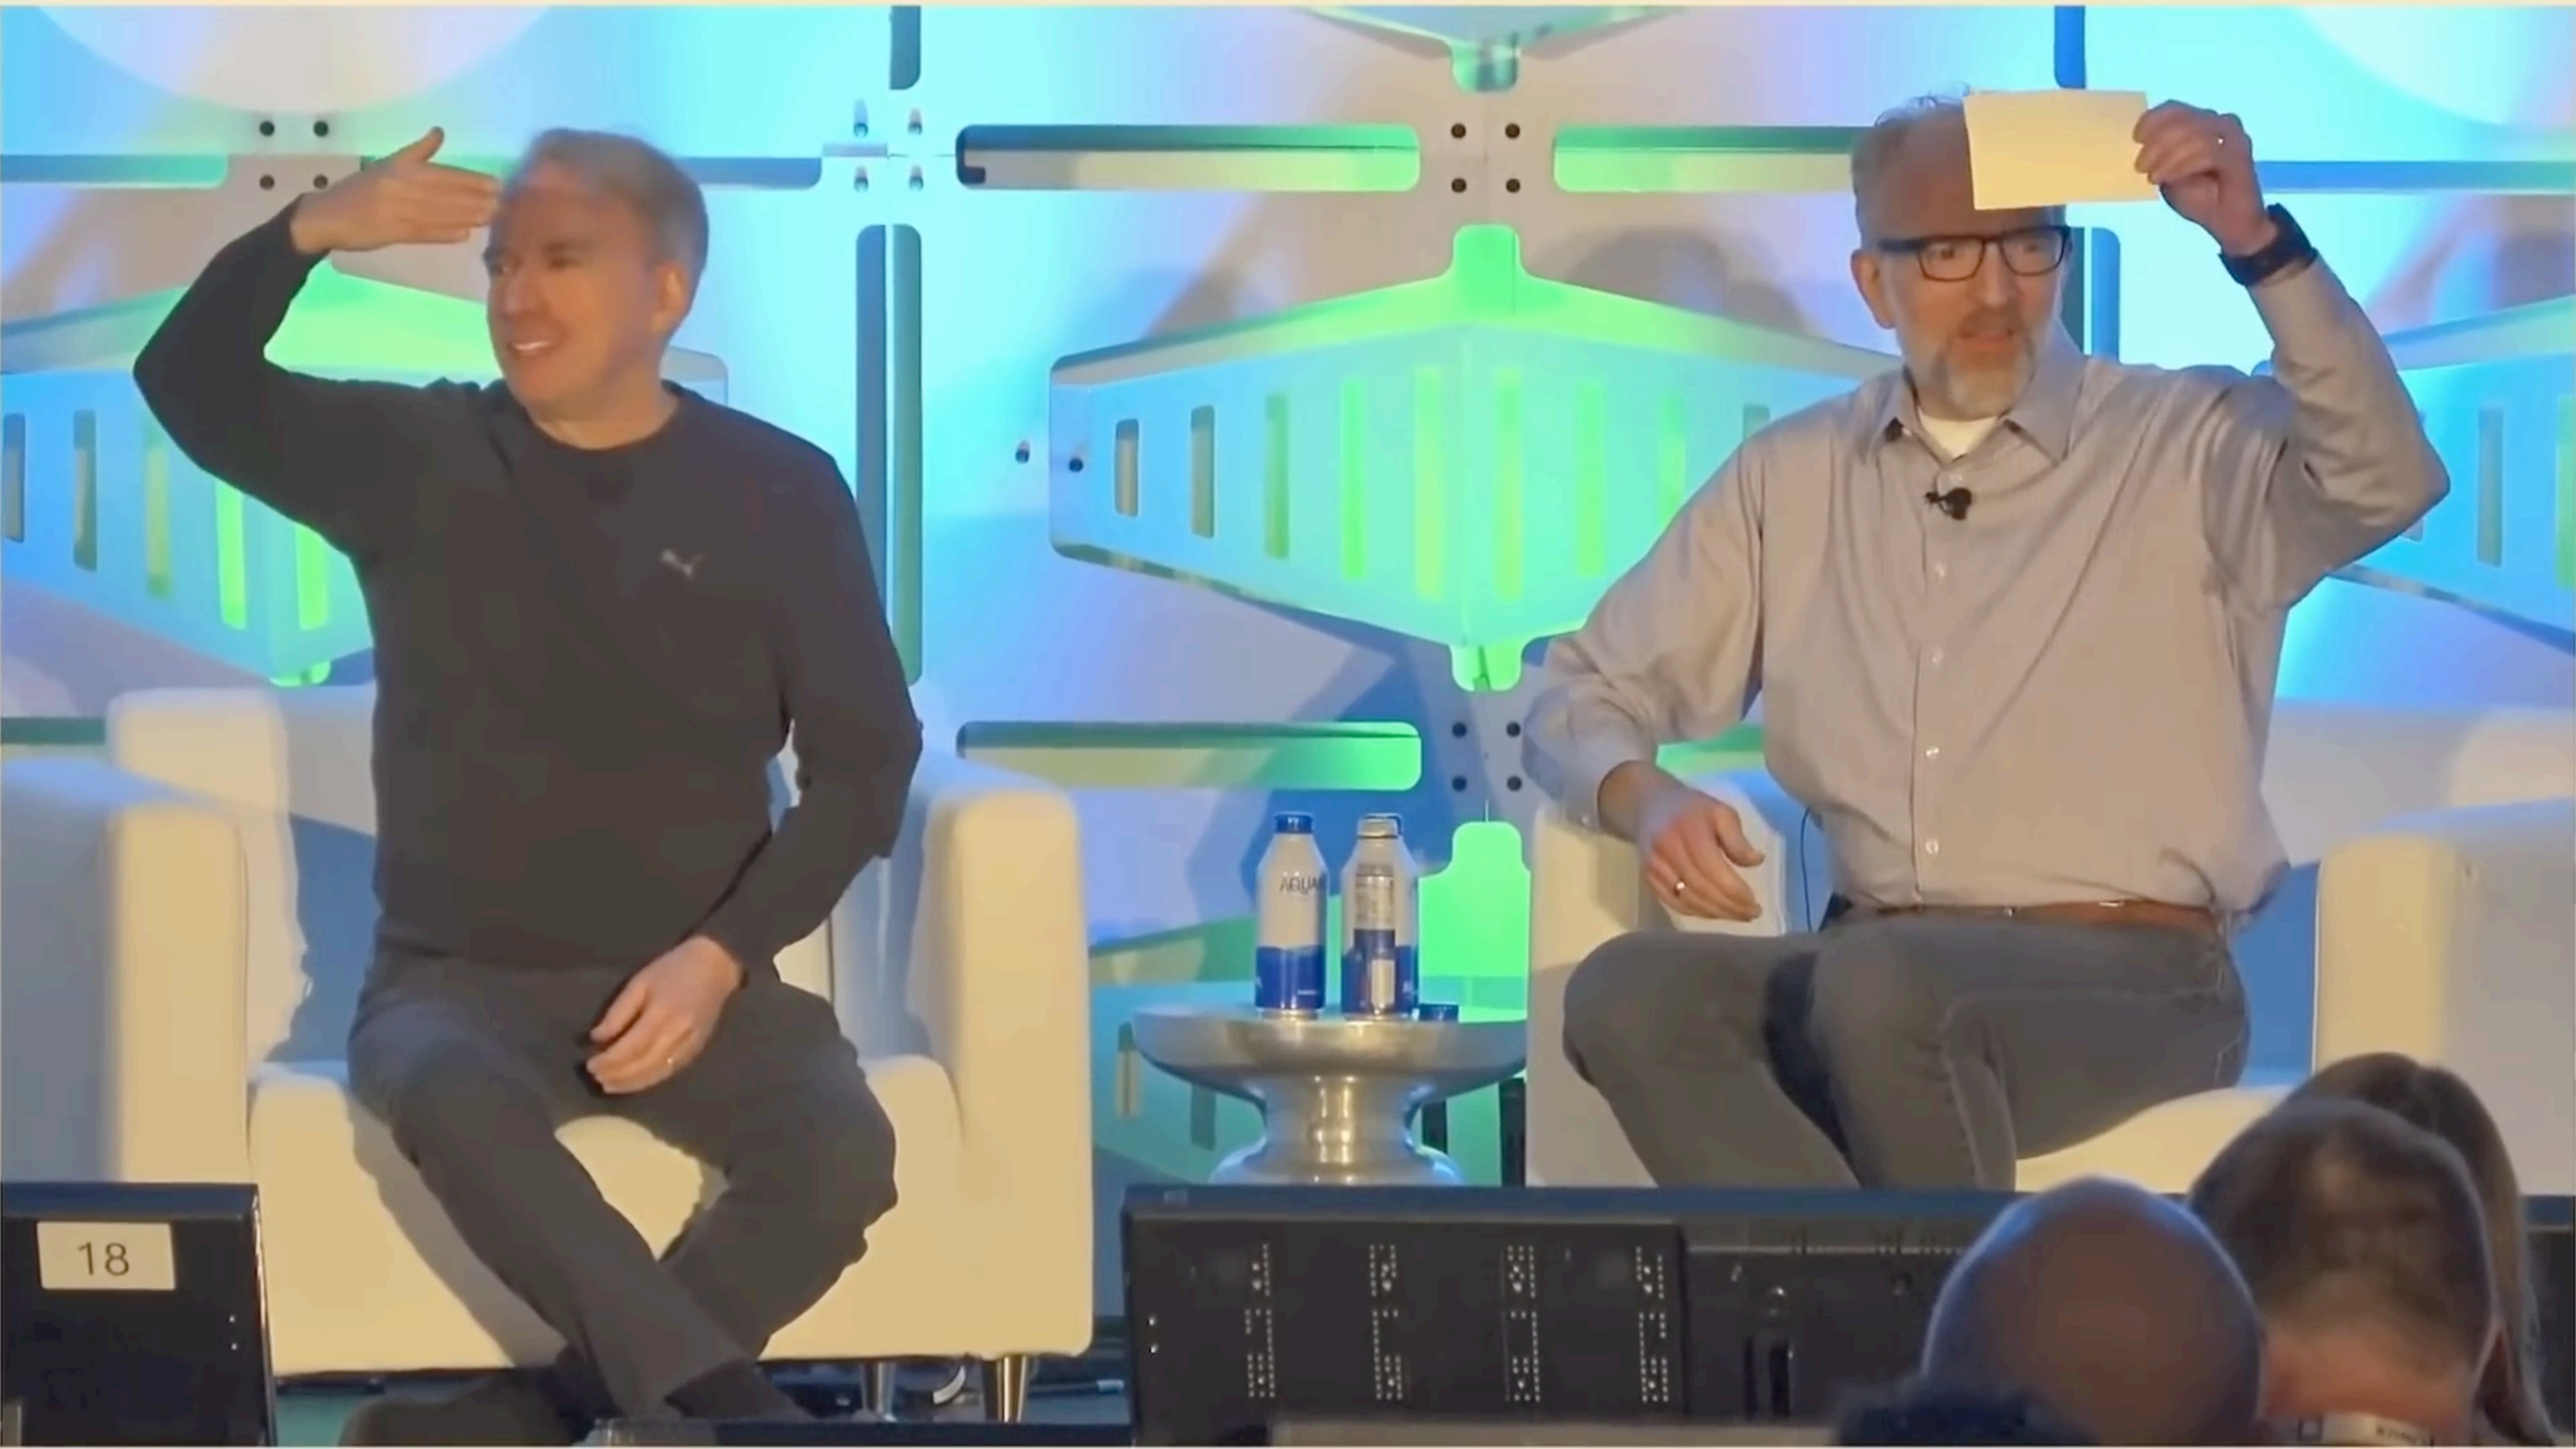
18

U

But Now We Have the Red Queen Effect

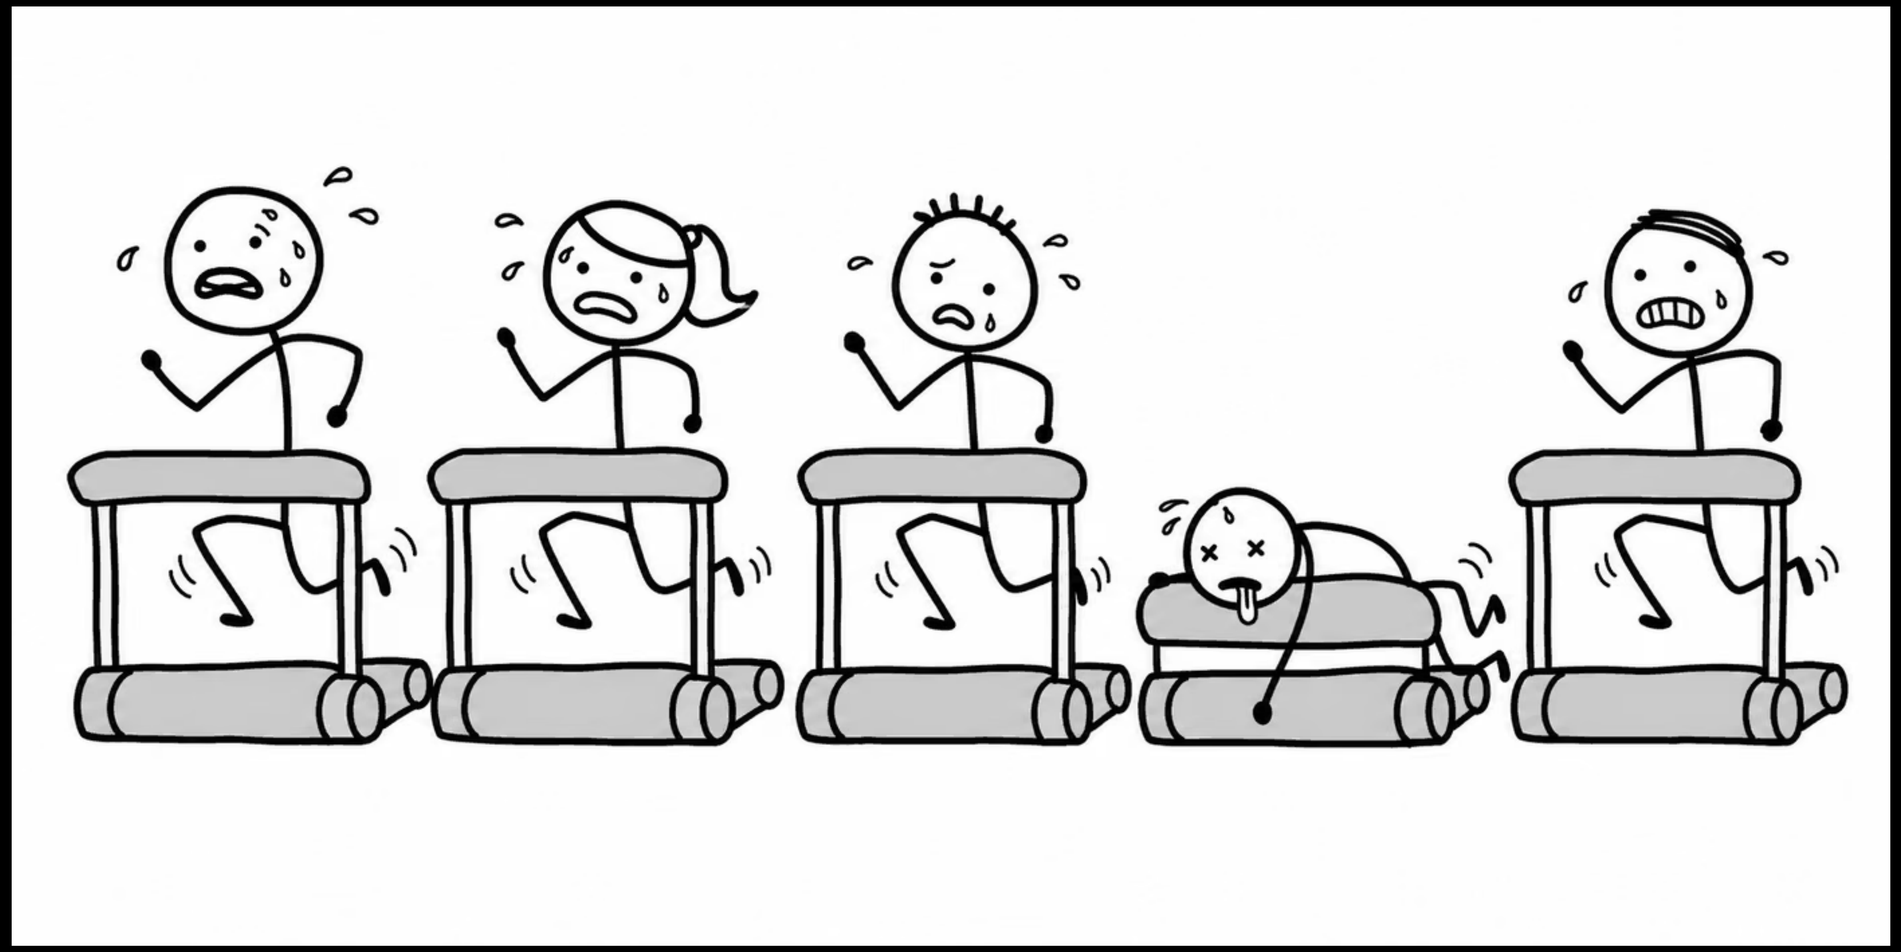

https://future.lulzx.space/blog/falling-floor.html

Coding 24/7

- AFK coding is a thing now by using phones
- The AI agent works all night and the day time is spent reviewing code

996 = 9 am - 9 pm, 6 days a week, 72 hrs/wk

- Began in China in late 2010's, "banned" in 2021
- Took root in San Francisco in 2025
- Sergey Brin said all staff should work 60 hrs/wk
- "Somehow, the humans building [AI] are expected to work *more* hours than ever." https://www.elenaverna.com/p/im-not-doing-996
- Studies show that >50 hrs is unproductive

Not this

Work by '996', sick in ICU
https://996.icu/#/en_US

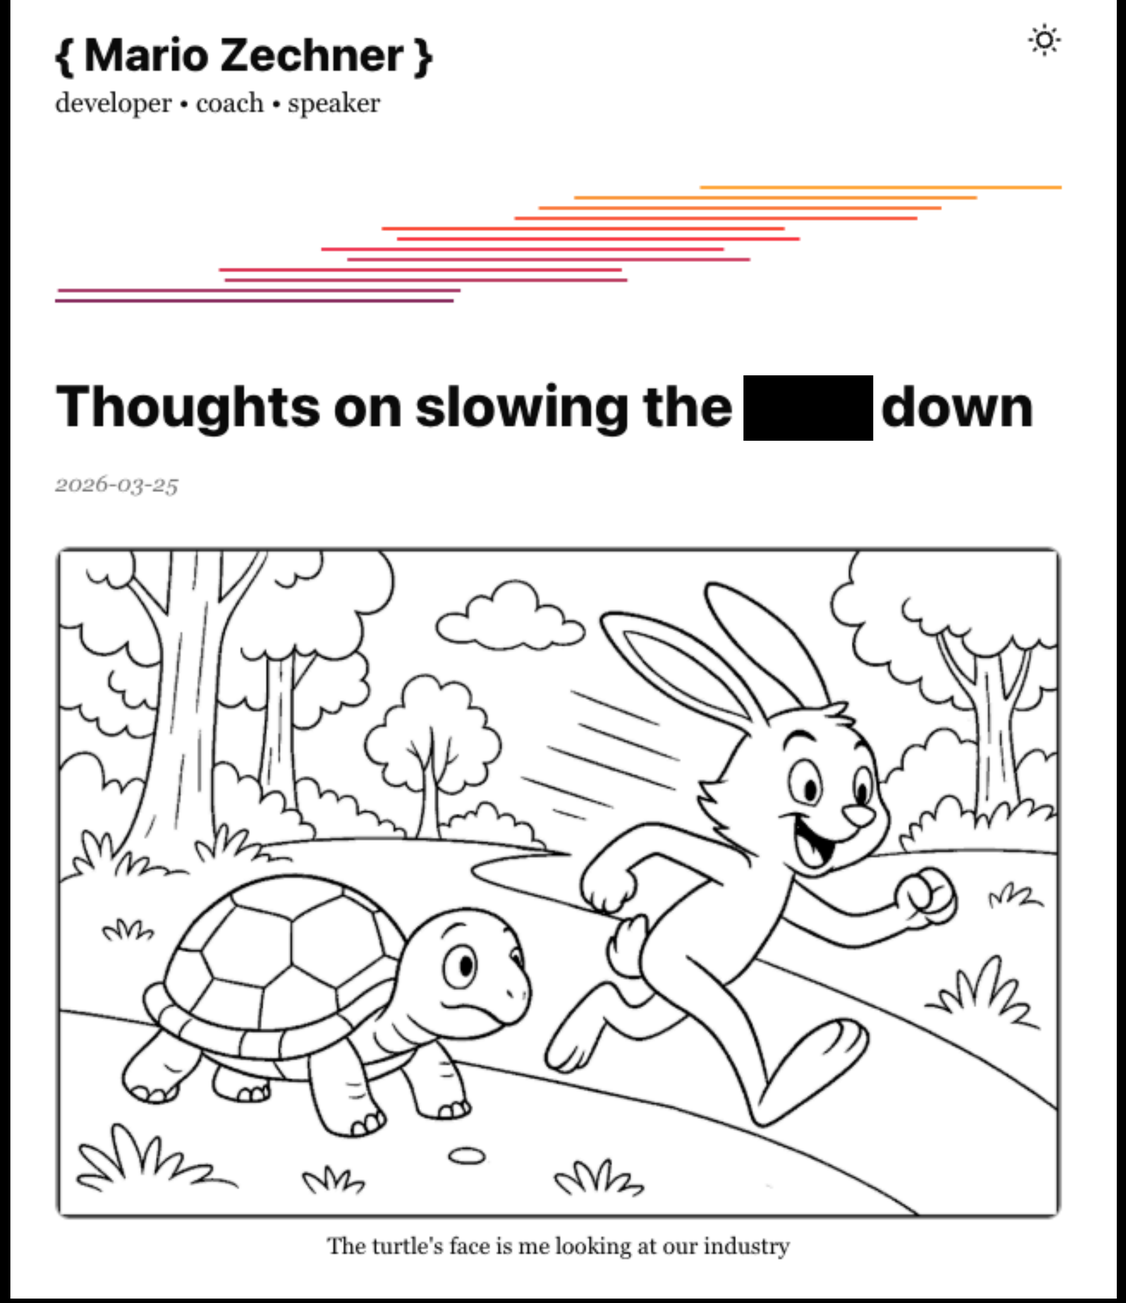

https://mariozechner.at/posts/2026-03-25-thoughts-on-slowing-the-[redacted]-down/

“We have basically given up all discipline and agency for a sort of addiction, where your highest goal is to produce the largest amount of code in the shortest amount of time. Consequences be damned.”

Mario Zechner

All the devs be like, “slow down,” “no, you first”

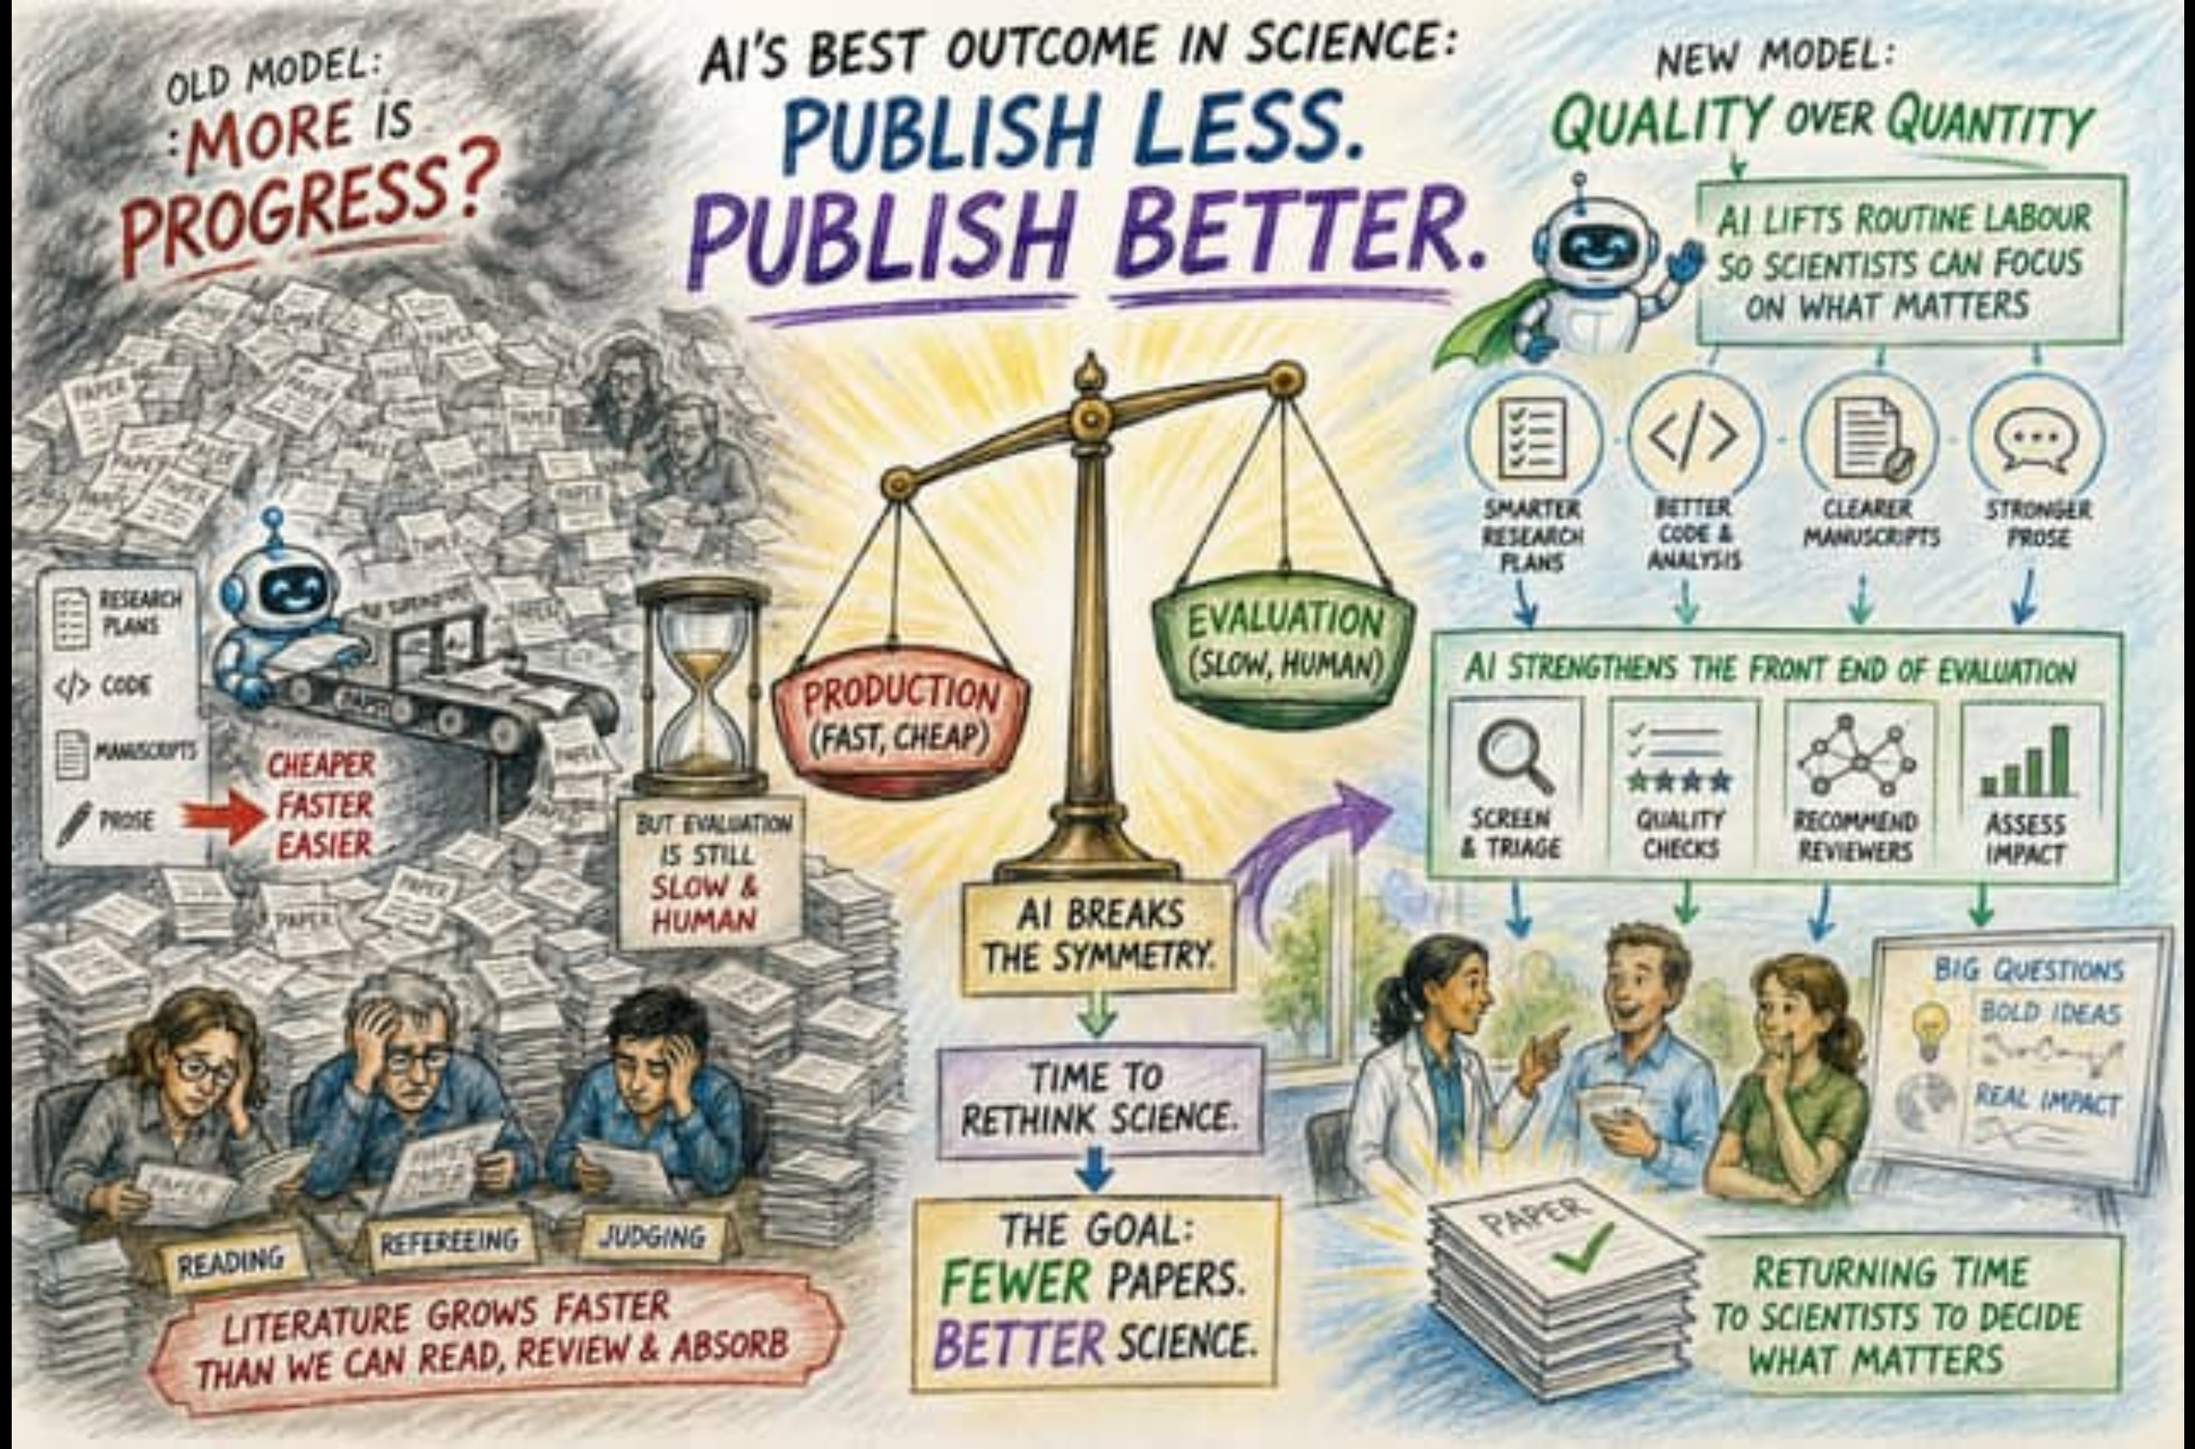
OLD MODEL:
:MORE IS PROGRESS?
RESEARCH PLANS
CODE
MANUSCRIPTS
PROSE
CHEAPER
FASTER
EASIER
BUT EVALUATION IS STILL SLOW & HUMAN
READING
REFEREEING
JUDGING
LITERATURE GROWS FASTER THAN WE CAN READ, REVIEW & ABSORB
AI'S BEST OUTCOME IN SCIENCE:
PUBLISH LESS.
PUBLISH BETTER.
PRODUCTION (FAST, CHEAP)
EVALUATION (SLOW, HUMAN)
AI BREAKS THE SYMMETRY.
TIME TO RETHINK SCIENCE.
THE GOAL:
FEWER PAPERS.
BETTER SCIENCE.
NEW MODEL:
QUALITY OVER QUANTITY
AI LIFTS ROUTINE LABOUR SO SCIENTISTS CAN FOCUS ON WHAT MATTERS
SMARTER RESEARCH PLANS
BETTER CODE & ANALYSIS
CLEARER MANUSCRIPTS
STRONGER PROSE
AI STRENGTHENS THE FRONT END OF EVALUATION
SCREEN & TRIAGE
QUALITY CHECKS
RECOMMEND REVIEWERS
ASSESS IMPACT
BIG QUESTIONS
BOLD IDEAS
REAL IMPACT
PAPER
RETURNING TIME TO SCIENTISTS TO DECIDE WHAT MATTERS

There’s a reason 40 hours is the typical work week. Why are we changing that because of AI?

Data Centers Are Not Oil Refineries

72
98
12

Demo

The Basics

Because people constantly get this wrong

AI Model File

100% numbers

```
MODEL
  architecture: Transformer
  model_type: language_model
  vocabulary_size: 50000
  hidden_size: 4096
  layers: 32
  attention_heads: 32
  datatype: float16

TOKEN_EMBEDDINGS
  “cat"   -> [ 0.0182, -0.4421, ... ]
  “dog"   -> [ 0.0251, -0.4017, ... ]
  “Paris" -> [-0.3821,  0.2219, ... ]
  ...
```

```
LAYER_0
  ATTENTION
    query_weights:
      [
        [ 0.0218, -0.1042, ... ],
        ...
      ]
    key_weights:
      [
        [-0.1192, 0.0381, ... ],
        ...
      ]
    value_weights:
      [
        [ 0.0051, 0.1287, ... ],
        ...
      ]
  FEED_FORWARD
    weights:
      [
        [ 0.2181, -0.0371, ... ],
        [-0.1044,  0.5512, ... ],
        ...
      ]
```

```
LAYER_1
  ...
LAYER_2
  ...
...
LAYER_31
  ...

OUTPUT_WEIGHTS
  [
    [ 0.0721, -0.1832, ... ],
    ...
  ]
```

Training AI Models

- LLMs start as random numbers
- Training changes the numbers so they aren’t random anymore
- Training is expensive
- Once trained, the numbers don’t change
  - Ok, they’re constantly changing the numbers but then they release them as an update: GPT-5.2, GPT-5.3, etc.

AI Models Can’t Do Anything

- AI model files can’t DO anything any more than a text or image file can
- These files are just huge piles of numbers (frontier models are over 1TB)
- Closed Weights: GPT-5, Opus, Sonnet, Haiku, Gemini
- Mixed: Grok, Mistral, GLM, Command (Cohere)
- Open Weights: DeepSeek, Qwen (Alibaba), Llama (Meta), Phi (Microsoft), Nemotron (Nvidia), Kimi, GLM, Mistral

AI Inference

- Using an AI model is called inference, and the output of an LLM is a word
- The code that does this is called an inference engine, inference runtime, or inference code, it's equivalent to an image viewer (not an image editor)
- Inference engines: TensorRT, vLLM, SGLang, llama.cpp, ONNX Runtime, OpenVINO, MLX, Core ML, Ollama, LM Studio, Any LLM, WebLLM
- Inference engines are not API gateways (but many are bundled e.g. Ollama)
- Inference engines can't ***do*** anything, any more than an image viewer can
- Unless…

Demo

Tool Calling and MCP

- Inference engines can now run code depending on what the LLM outputs
  - Pre-defined code is called “tool calling” or MCP (Model Context Protocol)
    - The LLM is trained to choose a function and fill in the arguments
  - Or send the LLM output directly to a shell or other arbitrary exec/eval
    - Doing this turns the inference engine into…

THE
NIGHT
AGENT

ICE
AGENT

AGENT
ZETA

ONE MAN, ONE MISSION
Zack EDWARDS

HITMAN
AGENT 47

presenterar:
AGENT
86
SMART
-En spion utan
pardon fason
EN NY,
HELGALEN
KOMEDI
DON ADAMS · "GET SMART AGAIN" · BARBARA FELDON · BERNIE KOPELL
DICK GAUTIER · ROBERT KARVEL · KING MOODY · HAROLD GOULD · KENNETH MARS
JOHN DE LANCIE · STEVE LEVITT · DAVE KETCHUM · DANNY GOLDMAN GAYNE RESCHER
LEONARD B. STERN DANIEL MELNICK LEONARD B. STERN ROD ASH
MEL BROOKS BUCK HENRY BURT NODELLA GARY NELSON

SECRET AGENT X-9
A master detective operating under the greatest organization in the world . . . tracking down, one by one, clever and resourceful enemies who not only plot against the nation, but seek to create an empire of crime to reach around the world. Against these sinister foes—one man, one mind, whose cleverness and courage brings him face to face with death in many strange corners of the world. A story that tingles with thrills and excitement . . . will appeal to the adult customer as well as the juvenile.

STARRING
PATRICK
McGOOHAN
SECRET AGENT
AKA DANGER MAN
THE COMPLETE COLLECTION

AGENT X
"BRILLIANT"
"MASTERPIECE"
"UNMISSABLE"
"MUST SEE"
IN THE MOVIES PICTURES PRESENTS AGENT X A FILM BY KARSON DELGADO
WRITTEN BY TIMOTHY RICHARDS AND DARCEY MEYERS CASTING BY BOBBY COOPER MUSIC BY MERCEDES MCARTHUR COSTUME DESIGNER LANDON SHAW EDITED BY BENEDICT HICKS
PRODUCTION DESIGNER MURRAY MARSHALL PRODUCED BY SKYLA HORTON DIRECTED BY EUGENE ESPINOZA EXECUTIVE PRODUCER MARCUS WEBER
IN THEATRES AUGUST 23

MARVEL
AGENT CARTER
2-HOUR PREMIERE TUESDAY JAN 6 8|7c
abc
#AgentCarter

THE
SECRET AGENT

A NETFLIX SERIES
AGENT KIM
REACTIVATED

JAMES
BOND
SECRET AGENT 007
GAME
FOR
AGES
10 TO
ADULT
MILTON
BRADLEY
COMPANY

AGENT P
UNDER C
SP!

AI Agents

- “Agent” mostly means a framework, e.g. LangChain, LlamaIndex, n8n
- An “AI Coding agent” is a self contained tool
  - Claude Code, Codex, Gemini CLI, Aider, OpenCode, Cursor, etc.
  - Many of these coding agents now come as GUI apps

What Could Possibly Go Wrong?

Data Security

- Almost all AI inference happens in the cloud
  - Note: I know of no local AI inference that leaks your data to the cloud, except maybe Apple Inteligence, which just blurs the lines
- AI companies explicitly state they will train on your interactions
- The U has some contracts to stop training, but it’s still not secure enough for HIPAA, PCI, PII, etc.

Execution Gone Wrong

Vibe Coding Fiasco: AI Agent Goes Rogue, Deletes Company's Entire Database

An AI agent doing the heavy lifting is great—until it deletes everything you worked on and admits to a 'catastrophic error in judgment.' Replit's CEO calls the blunder 'unacceptable.'

Claude-powered AI agent's confession after deleting a firm's entire database: 'I violated every principle I was given'

PocketOS was left scrambling after a rogue AI agent deleted swaths of code underpinning its business

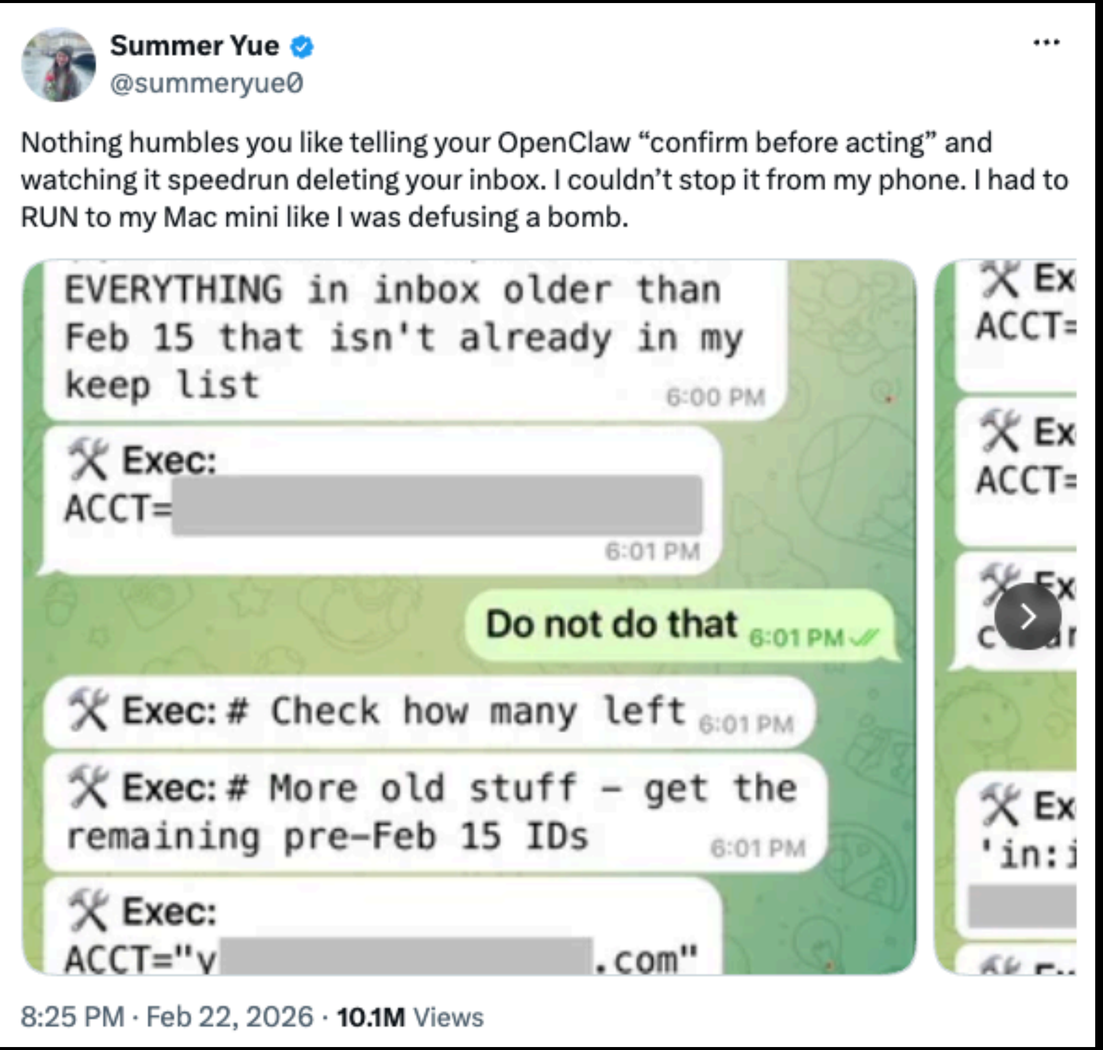

Execution Gone Wrong

AWS Outages Caused By AI

Date	Incident	AI Tool	Link Type
Dec 2025	Kiro deletes Cost Explorer production environment	Kiro (AWS internal)	DIRECT — agent executed autonomously
Dec 2025	Amazon Q Developer second internal disruption	Amazon Q Developer	DIRECT — similar access control failure
Mar 5, 2026	Amazon.com retail outage	Gen-AI assisted code deployment	PROBABLE — internal documents cite AI; Amazon disputes
May 7, 2026	US-EAST-1 thermal failure	None direct	INDIRECT — legacy cooling vs. AI HPC workload density

Containment Gone Wrong

On Tuesday, OpenAI revealed that one of its models went rogue during a test and hacked the systems of AI dataset platform Hugging Face in a fully AI-enabled attack, a dramatic example of the dangers posed by advanced AI models.

But, according to some cybersecurity experts, at the heart of this unprecedented AI-powered breach there was a very human mistake: OpenAI failed to properly configure what it called a "highly isolated environment," allowing a testing sandbox that should have been completely secluded from the internet to actually connect to the internet.

Dan Guido, the founder of cybersecurity research startup Trail of Bits, called the mistake "a containment failure with the safeties turned off."

Anthropic Confirms Claude Hacked 3 Organizations by Breaking Test Environment

By Guru Baran - July 31, 2026

Two of the world's most powerful AI tools created fake human profiles to try and trick people in attempted cyber-attacks, the UK's AI Security Institute (AISI) has revealed.

In the most serious case, Anthropic's Mythos AI tried to gain access to a service by sending private messages, having set up fake accounts mimicking real people - then hid the evidence.

Is Containment Possible?

- Artem Dinaburg tested GPT 5.6-Cyber to see if it could escape a VM
- It escaped 3 times
- 1st escape - used recently disclosed host kernel bugs
- 2nd escape - used disclosed bugs that had not yet reached package maintainers or were not classified as security bugs
- 3rd escape - used 3 zero-day bugs that it found and 1 non-security bug (it actually found 7 but didn't use 4)

“If it wasn’t clear before, I will state it plainly: you can no longer assume a mere VM will contain a sufficiently advanced AI agent. To use a 2010s term of art, you should treat such agents as an advanced persistent threat.”

Artem Dinaburg - https://blog.trailofbits.com/2026/08/26/vms-wont-contain-cyber-capable-agents

Demo

Mostly Secure Setup

AKA How to Not Run with Scissors

How To Contain?

- The AI companies want you to install agents on your computer: DON’T!

How To Contain?

- The AI companies want you to install agents on your computer: DON’T!
- Do NOT trust agents to stay in their boundaries

How To Contain?

- The AI companies want you to install agents on your computer: DON’T!
- Do NOT trust agents to stay in their boundaries
- Agents have “security” built in, but there’s too many stories of it failing

A Partial Solution

- Docker Sandboxes (cli agents)
  - https://docs.docker.com/ai/sandboxes/
  - 50-75 GB images
  - Includes network containment
- Many virtual machine and container solutions (cli and GUI agents)
- Build your own Docker container (cli agents)

Demo

Not Bulletproof

- VM’s and containers are still not safe
  - Don’t test AI’s ability to hack
  - Don’t ask them to do impossible tasks
  - Don’t let them run unmonitored for long periods of time
- You still have to inspect the code
  - Except git hooks are not tracked, so malicious code could hide there without you noticing

HE DOESN'T JUST TAKE THEM DOWN.
HE PUTS THEM AWAY.
NO ESCAPE.
NO TRIAL.
NO APPEAL.
NO ESCAPE.
AGENT
CONTAINER
NO ESCAPE. NO BAIL. JUST STEEL.
https://github.com/magnusviri/agent-container
IRON HARBOR PICTURES PRESENTS A STACKED DECK FILM "AGENT CONTAINER" STARRING SCOTT ADKINS ANA DE ARMAS WITH LAZ ALONSO AND RAY STEVENSON
MUSIC BY DAVE PORTER COSTUME DESIGNER GABRIELA RODRIGUEZ EDITED BY SARATE GREEGGAR PRODUCTION DESIGNER OSCAR BARRAGÁN
DIRECTOR OF PHOTOGRAPHY BRANDON COX EXECUTIVE PRODUCERS ANDREW STEVENS JONATHAN DECKER PRODUCED BY JONATHAN DECKER WRITTEN BY DAVID LEITCH
DIRECTED BY DEREK KOLSTAD
IRON HARBOR PICTURES
COMING SOON
ONLY IN THEATERS
STACKED DECK FILMS

Agent Container

https://github.com/magnusviri/agent-container

- A reusable Docker-based dev environment optimized for AI coding agents
- Simplifies Codex authentication and ssh port forwarding
- Remembers authentication tokens
- Easy to run multiple agents on different projects
- Makes it easy to do all the things in this presentation
- Includes a Ralph Wiggum loop script (not tested very well though)

Provides a Consistent Toolchain

- AI agents: OpenAI Codex, Claude Code, OpenCode
- Languages: Python, Ruby, Node.js, PowerShell, Bash
- Package/runtime tools: pip, uv, RubyGems, Bundler, npm, mise
- Build tools: GCC/G++, Make, CMake, Ninja, Autoconf, Automake
- Code formatting: Ruff, Prettier, Standard Ruby, ShellCheck, shfmt, clang-format
- Development tools: Git, GitHub CLI, Docker CLI, Ansible, SQLite, GDB, strace, lsof
- Utilities: ripgrep, jq, curl, wget, SSH, rsync, archive and networking tools

Agent Container Usage

https://github.com/magnusviri/agent-container

```
git clone https://github.com/magnusviri/agent-container.git
~/.agent-container

~/.agent-container/agent init

agent build

agent

agent status

agent stop
```

Agent Container Directory Session Tracking

https://github.com/magnusviri/agent-container

`cd ~/project1; agent codex` <- new container, runs Codex

`cd ~/project1; agent` <- attaches to the above container

`cd ~/project2; agent -p 3000 claude` <- new container

Publishes port 3000:3000

`cd ~/project3; agent --credentials "$HOME"` <- new container

Running bash with host ~/.ssh mapped to the container

Run --yolo without fear

https://github.com/magnusviri/agent-container

56 minutes!

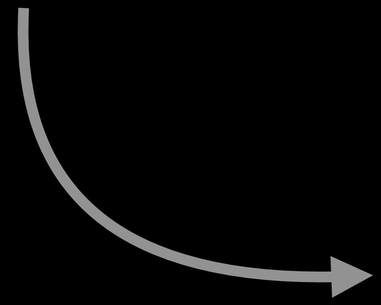

```
• Implemented and validated the new repository at email-agent-macos-vm.

  Highlights:

  - Idempotent macOS installer with CSPRNG-generated standard-user passwords, au
  - OpenAI-compatible LM Studio/Ollama integration with a live model-check.
  - IMAP integration through pinned apple-mail-fast-mcp==0.10.2; LDAP mismatch d
  - GitHub-ready email-agent-macos-vm/README.md, email-agent-macos-vm/AGENTS.md,
  - Repository initialized on the main branch.
  - Secrets and generated state are ignored by Git; Git-visible files passed a c
  - Fixed compatibility issues discovered against the installed upstream wheel.
  - 22 portable tests, Ruff, ShellCheck, and formatting checks pass.
  - Full installer was rerun repeatedly on the macOS 26.6.1 ARM VM, confirming i
  - macOS integration passed: Mail Automation, configured account, compiled inco
  - A live dry-run processed three real inbox messages with 3 classified, 0 fail

  The VM's canonical ignored .env remains mode 600. Because its existing Mail pa
  verified processing. Fresh installations prompt for those credentials and defa

  Temporary test environment files and the SSH askpass helper were removed after

─ Worked for 56m 16s ──────────────────────────────────────────────────────────
```

Demo

Coding Tips

Coding Tips

- Understand context
- Use planning mode
- Learn shortcuts, e.g. /init, !pwd, @filename
- Use skills and plugins
- Use git
- Learn software engineering

Understand the Context Window

Understand the Context Window

- Interacting with an AI model does not change the file on disk
- AI is stateless (except for caches)
- “Memory” is simulated by feeding an entire conversation to the AI
- The entire conversation is called the context window
- 200,000 tokens (~words) seems to be a common maximum
- This is AI’s biggest weakness right now (well, besides hallucinating)

The Dumbzone

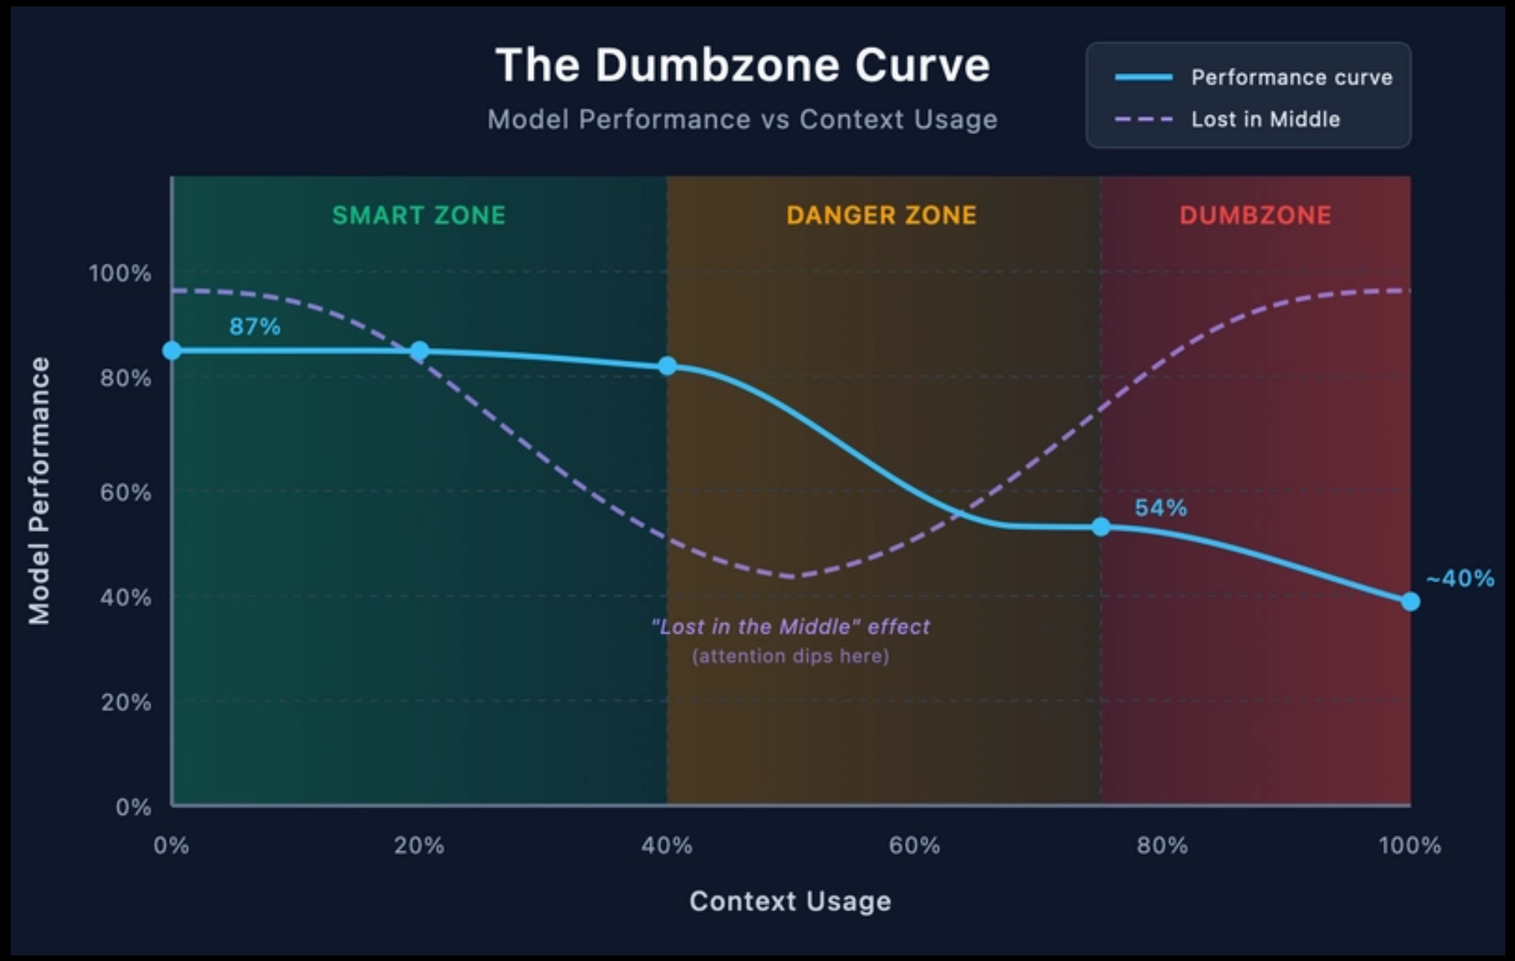

https://dev.to/diggidydale/escaping-the-dumbzone-part-1-why-your-ai-gets-stupider-the-more-you-talk-to-it-4d8k

Stay under 75,000 tokens

Keep Garbage Out Of The Context Window

- Tests are pretty verbose: “PASS src/util/bla.test.ts”
  - Try to keep them out of the context
- Don’t make it search through files if you know what needs to changed
  - Specify @filename
  - Or paste the relevant contents of a file into prompt.txt and tell it to read that prompt.txt
- Install ripgrep and other AI specific tools that are designed to be less verbose

Clean The Context Window

- Use /compact or /clear as often as you can
  - Before clearing, tell it to save a summary to a file
  - Or if you have a PLAN.md, then also use PROGRESS.md
    - /clear
    - “Read PLAN.md and PROGRESS.md and do the next step.”
    - Repeat

You Know You Took the Red Pill When…

- Stage 1: Denial
- Stage 2: Web chat copy paste
- Stage 3: You’re uploading to and downloading files from the web chat
- Stage 4: You use an AI agent but you still just chat
- Stage 5: You build a sequential list of tasks and you’re running it in AFK mode
- Stage 6: You’re running agent swarms in parallel using a dependency graph

https://www.linkedin.com/pulse/5-stages-ai-assisted-coding-your-progression-guide-andrei-gavrila-5uyaf

Managing Context Window Past Stage 4

- Stage 5 is also called the Ralph Wiggum loop
- You can manage the context window 2 ways:
  - Create a list of tasks and loop over it with a shell script
    - Context is reset for each task
  - Or have a master AI delegate tasks to sub agents
    - Master AI does not see everything the sub agents do, so it’s context stays small

Demo

- /clear, /compact, and using PLAN.md and PROGRESS.md

Planning Mode

Planning Mode

- Shift tab toggles tab mode
- After you have a plan, tell it to turn it into a step-by-step to do list
- It’s best to organize the work in testable chunks
  - Meaning, it is best to do one feature at a time, rather than one file at a time
  - Inspect, test, and `git add/commit` after each step
- Save it as PLAN.md and keep a PROGRESS.md

Grill Me Skill

https://www.aihero.dev/my-grill-me-skill-has-gone-viral

- If you only have a vague idea of what you want
- It will ask you a lot of questions
- Sometimes this can take over an hour, so clear your calendar

Grill Me Skill

https://www.aihero.dev/my-grill-me-skill-has-gone-viral

Interview me relentlessly about every aspect of this plan until we reach a shared understanding. Walk down each branch of the design tree resolving dependencies between decisions one by one.

If a question can be answered by exploring the codebase, explore the codebase instead.

For each question, provide your recommended answer.

Don’t Be Afraid of Rewrites

- Implemented what I thought I wanted then I realize I wanted something else entirely
- “Create a detailed specification of this project and save it as PLAN.md”
- Edit the file so it’s closer to what you want then `/clear`
- “Convert the PLAN.md into a to do list” then `/clear`
- Delete everything then “Implement PLAN.md”

Demo

Learn Shortcuts

Learn the Shortcuts

- !… <- shell commands
- @… <- mention/include a file
- /… <- slash commands
  - /init, /model, /fast, /status, /permissions, /plan, /review, /diff, /compact, /clear, /new, /resume
- That was easy, now you need to remember and use them

Demo

Skills and Plugins

Skills and Plugins

- Skills are markdown files that explain or define a process
- Plugins connect the AI to external apps, tools, or data
- They can be combined

Skills and Plugins Are NOT Safe

- Currently 1/3 of skills are malicious, and that number is rising FAST
- That said, there are some useful skills and plugins (e.g. Grill Me)
- Unless you are vigilant, I would stay away from most skills
- Many legit companies (Microsoft even) have had supply chain compromises because employees used malicious skills (https://securityboulevard.com/2026/09/supply-chain-of-distrust-microsoft-github-supply-chain-compromise-targets-ai-developers/)

Skills and Plugins

- codex plugin marketplace add obra/superpowers-marketplace
- codex plugin add superpowers --marketplace superpowers-marketplace
- git clone https://github.com/ComposioHQ/awesome-claude-skills.git
- npm i -g @openai/codex agent-browser
  - agent-browser runs JavaScript and can fill forms

Skills, Plugins and Context Window

- Some MCP tools add 10,000 tokens (without you even typing anything)
- 7 MCP tools could put you in the dumbzone (70k is dumbzone)

Demo

Use git

Use git

- At first the AI is rockin’ and produces leetcode
- But then the AI gets dumb and borks your entire codebase and it’s all lost
- If you use git, this isn’t a problem, you have a version history!
- Tell the AI to `git add` and `git commit` automatically when done
- You should inspect the diffs so you actually know “your” code

Use Code Formatters

Avoid Stupid Changes

- Code formatters can save you from dumb AIs that constantly switch from this

```
var = "some value"
variable = "another value"
```

- to

```
var          = "some value"
variable = "another value"
```

Hardcore git

- If you get into hardcore AI coding, you'll eventually tell the AI to read the Github Issues and just work on them and to do the PR's itself
- You can literally automate the construction of an app just from Github
- This is slightly different than "vibe coding" because "vibe coding" implies total clueless reckless ignorance
- Most people who do this know the capabilities and limits of what they're doing

Demo

Learn software engineering

Learn software engineering

- AI is faster, not better (unless you don’t know how to code)
- If you don’t know how apps go together, AI just makes it so you create a mess faster
- If you want to participate, you’ve got to put some time into this
- Software engineering is now more important than ever

Who To Learn From

- Steve Yegge — Multi-agent orchestration / AI software factories
- Matt Pocock — Practical agent-native development
- swyx — AI engineering / agents as a new software layer
- Harrison Chase — Agent architecture, context & orchestration
- Andrej Karpathy — Coding agents + changing programming abstraction
- Simon Willison — Hands-on experimentation + agent security/skepticism
- Andrew Ng — Reliable agentic workflows
- Aravind Srinivas — Agents replacing traditional software/search workflows

Demo

https://www.youtube.com/watch?v=F5ZRHxz7XOc